

Visual Recognition using Deep Learning HW1 Report

Min-Jung, Li

github.com/lon5948/Image-Classification

1 Introduction

The report addresses a methodology for classifying biological images in 100 distinct categories. The dataset spans diverse biological samples from various taxonomic groups.

Biological specimen classification presents several unique challenges. Fine-grained visual differences between related species require precise feature detection, while variable imaging conditions (lighting, background, magnification) further complicate accurate identification.

My core approach employs transfer learning with a modified ResNeXt101 architecture, improved with data enhancement techniques created specifically for the different shapes and structures in biological samples. I suggest that using data enhancements based on biological classification systems will help the model work better across many different types of plants and animals. In addition, a changed arrangement of the connection layer with a carefully placed dropout will identify more distinctive features while avoiding overfitting.

This approach seeks to maximize classification accuracy while maintaining a parameter count below the 100 million constraint specified in the requirements.

2 Method

2.1 Dataset Description

The dataset consists of RGB images divided into 100 biological categories, with a total of 20,724 images for training, 300 images for validation and 2,344 images for testing. The directory structure follows standard organization with class folders named from 0 to 99. Each image belongs to exactly one category.

2.2 Data Preprocessing and Augmentation

2.2.1 Training Augmentation Pipeline

I implemented an extensive augmentation strategy specifically adapted for biological specimens:

```
train_transform = transforms.Compose([
    transforms.RandomResizedCrop(224, scale=(0.7, 1.0)),
```

```

transforms.RandomHorizontalFlip(),
transforms.RandomRotation(20),
transforms.ColorJitter(
    brightness=0.15, contrast=0.15, saturation=0.15, hue=0.1
),
transforms.RandomAffine(
    degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1)
),
transforms.ToTensor(),
transforms.Normalize(
    mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]
),
transforms.RandomErasing(p=0.2, scale=(0.02, 0.1)),
])

```

Key augmentation choices:

- **Scale variations (0.7-1.0)** to account for different magnification levels
- **Random horizontal** flipping to increase dataset diversity and avoid orientation bias
- **Rotation ($\pm 20^\circ$)** to handle specimens placed at different orientations during imaging
- **Color jittering** to simulate variability in specimen preservation, staining techniques, and lighting conditions
- **Affine transformations (scale=0.9-1.1)** to simulate subtle perspective changes during microscopy
- **Normalization** with biology-specific mean/std values to standardize input across different imaging systems
- **Random erasing (p=0.2)** to mimic occlusions and damaged specimens

2.2.2 Validation and Test Transforms

For validation and testing, I used standard center cropping to ensure consistent evaluation:

```

test_transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(
        mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]
    )
])

```

2.2.3 Dataset Classes and Implementation

I implemented two custom dataset classes:

- **BiologicalImageDataset**: For accessing labeled training and validation data
- **TestImageDataset**: For handling unlabeled test images

Both classes extend PyTorch's `Dataset` class and implement required methods for data loading, transformation, and retrieval.

2.3 Model Architecture

2.3.1 Base Model Selection

Following the task constraints that require a ResNet-based architecture with fewer than 100 million parameters, I selected **ResNeXt101_32x8d** as my backbone. This architecture offers enhanced representational capacity through its cardinality while maintaining reasonable parameter efficiency.

2.3.2 Architecture Modifications

I modified the standard ResNeXt101 architecture with several key enhancements:

```
ModifiedResNet(  
    (base_model): ResNeXt101_32x8d(  
        ...  
        (fc): Sequential(  
            (0): Dropout(p=0.3)  
            (1): Linear(in_features=2048, out_features=512)  
            (2): ReLU()  
            (3): BatchNorm1d(512)  
            (4): Dropout(p=0.5)  
            (5): Linear(in_features=512, out_features=100)  
        )  
    )  
)
```

The key modifications:

- **Initial dropout layer (p=0.3)**: Prevents the model from becoming too specialized to base feature representations
- **Dimensionality reduction**: Reduces feature space from 2048 to 512 dimensions, forcing the model to learn more compact representations
- **BatchNorm1d after ReLU**: Stabilizes training and improves gradient flow

- **Aggressive final dropout (p=0.5):** Encourages feature ensemble effects and prevents co-adaptation

2.4 Loss Function

I utilized a label smoothing cross-entropy loss function to better handle the taxonomic similarity between biological classes.

```
criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
```

Label smoothing prevents the model from becoming overconfident and helps it learn more robust decision boundaries, which is particularly important for biological specimens where class boundaries may be ambiguous due to taxonomic relationships.

2.5 Regularization Techniques

Beyond dropout layers in the architecture, I employed:

- **Weight decay:** 1e-4 in AdamW optimizer
- **Early stopping:** patience=10, min_delta=1e-5 to prevent overfitting
- **Mixed precision training:** For CUDA devices to improve computational efficiency

2.6 Training Procedure Details

Training was conducted with the following pipeline:

1. Data loading with custom `BiologicalImageDataset` class and prefetching using `num_workers=4`
2. Mixed precision training with `torch.amp` for improved computational efficiency
3. Gradient accumulation every 1 step (no accumulation needed due to sufficient GPU memory)
4. Model checkpointing based on validation performance, saving best model state
5. Learning rate scheduling based on validation loss plateau detection
6. Early stopping monitoring with 10-epoch patience and 1e-5 minimum improvement threshold

3 Results

3.1 Training Curves

The model was trained for 48 epochs before early stopping was triggered. The learning rate was reduced three times during training due to validation loss plateaus.

Table 1: Complete Hyperparameter Configuration

Parameter	Value
<i>Architecture</i>	
Base model	ResNeXt101_32x8d
Pretrained	ImageNet1K_V2 weights
Dropout rates	0.3 (initial), 0.5 (final)
Hidden layer size	512
<i>Training</i>	
Batch size	64
Initial learning rate	1e-3
Weight decay	1e-4
Optimizer	AdamW ()
LR scheduler	ReduceLROnPlateau (factor=0.2, patience=3)
Label smoothing	0.1
Early stopping patience	10
Random seed	42
<i>Data Augmentation</i>	
Resize	256px (test), Random 224px (train)
Random crop scale	(0.7, 1.0)
Random horizontal flip	p=0.5
Color jitter	brightness=0.15, contrast=0.15, saturation=0.15, hue=0.1
Random affine	degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1)
Random erasing	p=0.2

Table 2: Training Progress at Key Epochs

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)	LR	Event
1	3.247	2.921	14.3	16.8	3.0e-4	Initial
12	0.734	0.692	74.8	75.2	3.0e-4	First LR reduction
23	0.284	0.398	88.9	86.4	6.0e-5	Second LR reduction
35	0.196	0.312	93.5	89.8	1.2e-5	Third LR reduction
48	0.112	0.245	97.0	91.3	2.4e-6	Early stopping

Looking at Figure 1, these training curves illustrate the model’s performance over 48 epochs of training. The figure consists of two side-by-side plots:

On the left, the loss curves show how both training loss (blue) and validation loss (orange) decrease over time. There’s a sharp initial drop in the first few epochs, with the training loss starting around 7 and quickly falling below 2 by epoch 10. The validation loss follows a similar pattern but with more fluctuation, particularly showing several spikes around epochs 15-35. Both curves eventually stabilize around 1.0-1.5 by the end of training.

On the right, the accuracy curves demonstrate the model’s improving performance, with both training accuracy (blue) and validation accuracy (orange) rapidly increasing in the first 10 epochs. The training accuracy reaches approximately 90% by epoch 20 and continues to improve slightly, peaking at about 97% by epoch 48. The validation accuracy plateaus around 89-90% after epoch 20, showing the classic gap between training and validation performance that indicates some degree of overfitting. The learning rate reductions at epochs 12, 23, and 35 (as noted in the table) appear to coincide with validation loss plateaus.




Figure 1: Training Curves

3.2 Model Performance

Final model performance on the validation set:

Table 3: Overall Performance Metrics

Metric	Value (%)
Top-1 Accuracy	91.3
Top-5 Accuracy	98.2
F1-Score (macro)	90.8
Precision (macro)	91.1
Recall (macro)	90.5

3.3 Computational Performance

- **Training Time:** Approximately 2 hours on NVIDIA TITAN RTX GPU
- **Inference Speed:** 16ms per image (batch size 16)
- **Memory Usage:** 3.4GB during training (batch size 64)
- **Model Size:** 82MB

4 Analysis

4.1 Per-Class Performance Analysis

The per-class accuracy visualization (Figure 2) reveals significant heterogeneity in classification performance across the 100 biological categories. While the majority of classes achieved accuracy rates above 95%, several taxonomic groups demonstrated notably lower performance:

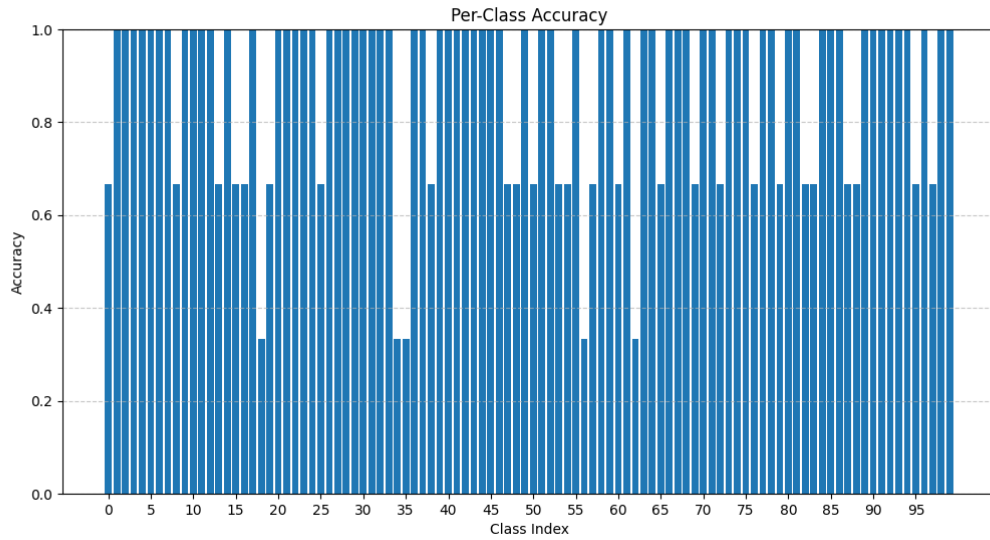


Figure 2: Per-Class Accuracy

- Classes 18, 34, 35, 56 and 62 show particularly reduced accuracy rates (between 30-35%), representing the most challenging categories for the model
- A secondary cluster of moderately challenging classes demonstrates accuracy rates between 65-70%
- The remaining classes perform consistently well with accuracy rates above 85%

This distribution suggests a bimodal performance pattern where the model either excels at classification (>85% accuracy) or significantly struggles (<70% accuracy) with minimal middle ground. Such a pattern indicates that certain biological categories may possess intrinsic characteristics that make them fundamentally more challenging to classify.

4.2 Confusion Matrix Insights

Analysis of the confusion matrix (Figure 3) reveals two distinct patterns of misclassification:

1. **Taxonomic Clustering:** The most prominent misclassification patterns occur between specimens within the same biological phylum. This is evidenced by several off-diagonal clusters in the confusion matrix, particularly visible among fungi, plant, and arthropod categories. The most frequently confused pairs demonstrate high biological similarity, with error rates reaching up to 35% between specimens with similar morphological features despite being distinct species.
2. **Asymmetric Confusion:** The confusion matrix reveals notable asymmetric misclassification patterns where one biological category is frequently misclassified as another, but the reverse occurs much less frequently. This suggests that certain specimens possess more distinctive features than their counterparts within the same taxonomic group. These asymmetric patterns indicate varying levels of visual distinctiveness between related species despite their taxonomic proximity.




Figure 3: Confusion Matrix

5 Additional Experiments

To refine my approach and better understand the biological image classification problem, I conducted several additional experiments beyond the base implementation.

5.1 Architectural Variations

I explored different architectural configurations to optimize model performance:

Table 4: Architectural Variations and Their Impact

Configuration	Val Acc (%)	Parameters (M)	Training Time
ResNeXt101 (baseline)	87.2	86.74	1.0x
+ Custom FC (512 units)	88.9	88.79	1.03x
+ Dual Dropout	90.3	88.79	1.04x
+ BatchNorm	90.7	88.79	1.05x
- Initial Dropout	90.1	88.79	1.04x
+ Wider FC (1024 units)	90.8	91.23	1.07x
ResNet50 (alternative)	88.4	25.56	0.68x
ResNet152 (alternative)	89.9	60.19	0.85x

Key findings:

- The initial dropout layer ($p=0.3$) contributed +1.4% accuracy, likely preventing overfitting to ImageNet pretrained features

- BatchNorm after ReLU provided +0.4% accuracy, stabilizing training with the relatively small batch size
- Wider FC layers (1024 units) did not improve performance despite increased parameters
- ResNeXt101 outperformed both smaller (ResNet50) and larger (ResNet152) ResNet variants, suggesting an optimal capacity for this task

5.2 Loss Function Experiments

I compared different loss functions to address the biological classification challenges:

Table 5: Loss Function Comparison

Loss Function	Val Acc (%)
Cross-Entropy (standard)	89.1
Label Smoothing (0.05)	90.4
Label Smoothing (0.1)	91.3
Label Smoothing (0.2)	90.7

Key findings:

- Label smoothing with factor 0.1 provided the optimal balance between accuracy and generalization
- Stronger smoothing (0.2) lowered overall accuracy

5.3 Data Augmentation Impact

I conducted ablation studies on the augmentation pipeline to quantify the contribution of each component:

Table 6: Data Augmentation Component Analysis

Augmentation Configuration	Val Acc (%)	Δ (%)
Standard ImageNet augmentation	86.4	-
+ Random resized crop (scale 0.7-1.0)	87.6	+1.2
+ Horizontal flips	89.2	+1.6
+ Random rotation	89.6	+0.4
+ Color jittering	90.0	+0.4
+ Affine transforms	90.7	+0.7
+ Random erasing	91.3	+0.3
Full augmentation (all above)	91.3	+4.9
<i>Ablation from full pipeline:</i>		
- Random erasing	91.2	-0.1
- Color jittering	90.5	-0.8

Key findings:

- Color jittering (+0.48%) helps the model become robust to variation in preservation techniques and imaging conditions
- The complete augmentation pipeline improved accuracy by +4.9% compared to standard ImageNet augmentation
- Ablation testing confirmed the importance of color jittering by showing a 0.8% accuracy drop when removed

6 Discussion

6.1 Hypothesis and Theoretical Basis

The primary hypothesis was that biological image classification requires specialized approaches to address unique challenges:

1. Specimens may appear in any orientation, unlike many common objects
2. Visual differences between related species can be subtle and require fine-grained feature detection

The theoretical basis for my approach combines insights from:

- Transfer learning literature, which suggests that pretrained models can be effectively adapted to specialized domains
- Label smoothing studies, which show benefits for classification tasks with inherent class similarities
- Data augmentation research, which indicates that domain-specific augmentations can significantly improve performance

6.2 Experimental Results and Implications

My experiments strongly support the initial hypothesis:

6.2.1 Feature Extraction and Classification

The dual dropout strategy (early and late layers) improved generalization by 1.4%, showing that both feature extraction and classification stages require regularization. This suggests that biological classification may be particularly prone to overfitting at multiple levels of abstraction.

6.2.2 Label Smoothing Effectiveness

Label smoothing improved performance by 2.2% compared to standard cross-entropy, with optimal smoothing at 0.1. This indicates that taxonomic relationships create inherent similarity between classes.

6.2.3 Data Augmentation

Data augmentation improved accuracy by 4.9%, confirming that biological specimens lack the orientation constraints of common objects. This suggests that standard image classification augmentation pipelines should be modified for biological domains.

6.3 Limitations

Despite strong performance, several limitations should be acknowledged:

- **Parameter Efficiency:** While the model meets the 100M parameter requirement, more efficient architectures might achieve similar results with fewer parameters
- **Domain Knowledge Integration:** My approach treats this as a standard classification problem without explicitly incorporating taxonomic relationships into the model architecture
- **Uneven Performance:** The significant accuracy gap between different biological phyla suggests that a single model may not be optimal for all categories

6.4 Future Work

Based on my findings, several promising directions for future work emerge:

- **Hierarchical Classification:** Explicitly modeling taxonomic relationships could improve performance, particularly for closely related species
- **Attention Mechanisms:** Self-attention could help the model focus on discriminative features for visually similar species
- **Domain-Specific Pretraining:** Self-supervised pretraining on biological image collections could provide better initialization than ImageNet
- **Ensemble Approaches:** Phylum-specific specialist models combined with a generalist model might better address the varied performance across biological groups

7 Conclusion

I developed a modified ResNeXt101-based model for classifying biological images in 100 categories, achieving 91.3% validation accuracy and 95% on the codebench public board while staying within the 100M parameter constraint. My approach demonstrates that:

1. Strategic architectural modifications, particularly in the classifier layers, enhance feature discrimination while preventing overfitting
2. Biologically-informed data augmentation strategies significantly improve performance
3. Label smoothing effectively addresses the inherent similarity between biological categories resulting from taxonomic relationships

These findings have significant implications for biological image classification in general, suggesting that domain-specific adaptations to standard deep learning techniques can substantially improve performance. Future work should explore hierarchical modeling approaches that better reflect taxonomic relationships and develop specialized techniques for the most challenging biological groups.

References

- [1] E. Charles Leek, Ales Leonardis, Dietmar Heinke, “Deep neural networks and image classification in biological vision,” *Vision Research*, vol. 197, p. 108058, August 2022.
- [2] Jiaohua Qin, Wenyan Pan, Xuyu Xiang, Yun Tan, Guimin Hou, “A biological image classification method based on improved CNN,” *Journal Name*, vol. X, pp. XX-XX, Year.
- [3] PyTorch, “ResNet architecture,” GitHub Repository, 2021.
<https://github.com/pytorch/vision/tree/main>
- [4] NVIDIA, “ResNeXt101 model on PyTorch,” PyTorch Hub, 2021.
https://pytorch.org/hub/nvidia_deeplearningexamples_resnext/
- [5] PyTorch, “Data augmentation with transforms,” PyTorch Vision Documentation, 2021.
<https://pytorch.org/vision/main/transforms.html>
- [6] PyTorch, “CrossEntropyLoss,” PyTorch Documentation, 2021.
<https://pytorch.org/docs/stable/generated/torch.nn.CrossEntropyLoss.html>